

Vasudev Menon

919-798-7081 | vmenon2@ncsu.edu | vmenon04.github.io | linkedin.com/in/vmenon04 | github.com/vmenon04

EDUCATION

North Carolina State University	Raleigh, NC
<i>B.S. in Computer Science and B.S. in Mathematics</i> Cum. GPA: 4.0/4.0	Expected Spring 2026
<ul style="list-style-type: none">Coursework: Deep Learning, Artificial Intelligence, Data Science, Linear Algebra, Data Structures and Algorithms, Operating Systems, Combinatorics, Real Analysis, Differential Equations, Discrete MathematicsHonors and Awards: Phi Beta Kappa, Dean's List All Semesters	
Eötvös Lorand University	Budapest, Hungary
<i>Budapest Semesters in Mathematics (BSM)</i>	Fall 2025
<ul style="list-style-type: none">Coursework: Deep Learning, Graph Theory, Advanced Combinatorics (Hypergraph Theory)	

EXPERIENCE

Fidelity Investments	Durham, NC
<i>Software/Machine Learning Intern - Summer 2025</i>	June 2025 – August 2025
<ul style="list-style-type: none">Engineered RAG system to automate translation of pension plan documentation into executable Java codeUsed AWS Bedrock, LangChain, HuggingFace, and Boto3 to build scalable ML infrastructureDesigned Neo4j graph database to encode pension rule relationships and dependencies for code generationOptimized LLM accuracy through prompt engineering, validation pipelines, and context management techniquesDeployed end-to-end Streamlit application integrating RAG pipeline, graph database, and code generation workflows	
Fidelity Investments	Durham, NC
<i>Software Engineer Intern - Summer 2024</i>	June 2024 – August 2024
<ul style="list-style-type: none">Engineered an enterprise assessment system used by 4,000+ company leads to evaluate application reliabilityImplemented a dashboard integrated with Power Platform APIBuilt and deployed Backstage developer portal plugins in Node.js and TypeScript, enhancing internal toolingUtilized Docker for DevDocs integration to provide documentation for users	
North Carolina State University	Raleigh, NC
<i>NSF-Funded Deep Learning Research Intern</i>	May 2023 – August 2023
<ul style="list-style-type: none">Automated deep learning experiments via shell scripting and R, analyzing the performance of batch normalizationTrained/evaluated neural nets with statistical analysis and metric visualizations (precision, recall, time)Contributed to NSF-funded research through experiment design and literature synthesis	
Rice University	Remote
<i>Data Science Intern (NSF Project)</i>	June 2021 – Feb. 2022
<ul style="list-style-type: none">Conducted research on an NSF-funded project to analyze and extrapolate COVID-19 sequence data using RUtilized numerous state-of-the-art statistical algorithms to examine DNA sequence dataAnalyzed 100,000+ COVID-19 strains using statistical models, uncovering key genomic differencesPresented research at the Junior Science and Humanities Symposium	

PROJECTS

Mechanistic Interpretability of Transformers on Simple Games <i>PyTorch, NumPy, HuggingFace</i>	
<ul style="list-style-type: none">Conducted interpretability research comparing hardcoded vs learned transformer models on controlled game tasksDesigned novel experimental framework for enabling direct comparison of hand-engineered and gradient-learned solutionsAnalyzed transformer internals to understand how neural networks learn algorithmic reasoning	
Differential Privacy in Federated Learning (FL) on MIMIC-IV <i>AI/ML, Python, PyTorch, Flower, Big Data</i>	
<ul style="list-style-type: none">Implemented Distributed Differential Privacy (DDP) in FL using Flower to protect sensitive healthcare dataSimulated heterogeneous clients using racial groupings from the MIMIC-IV ICU dataset to reflect real-world skewAnalyzed trade-offs between privacy and model accuracy, with Gaussian noise and clipping techniquesOvercame large-scale (100GB+) data processing challenges with SQL and Power QueryCollaborated with Dr. Kotevska (ORNL) and Dr. Haider (NCNU); authored and presented findings	
Rapster: Full-Stack Music Sharing Platform <i>AI/ML, FastAPI, Supabase, Docker, Essentia, Cloudflare R2</i>	
<ul style="list-style-type: none">Developed a full-stack app to upload and analyze audio tracks with ML-powered feature extraction (BPM, key, MFCCs)Deployed ML workloads using Essentia containerized with Docker and orchestrated via Docker ComposeStored user files in Cloudflare R2 and metadata in Supabase, using signed URLs for access control	

TECHNICAL SKILLS

Programming: Python, Java, C/C++, R, SQL
AI/ML: PyTorch, TensorFlow, Keras, Scikit-learn, HuggingFace, Federated Learning, Differential Privacy
Systems & Tools: Docker, AWS, Boto3, Git, CI/CD, Power Platform, Linux, REST APIs
Frameworks: FastAPI, Flask, Django, Flower, Streamlit
Other: Data Visualization (Matplotlib, Pandas), L ^A T _E X